

# LING001\_Final\_Project

December 20, 2018

## 1 LING001 Final Project

### 1.1 Introduction

The Vector Space Model for information retrieval is an interesting linguistic analysis technique. The paramount issue with language processing is representing words mathematically to do relevant computations. So, the Vector Space Model was developed to allow words to be represented as vectors! Using the model, we can find the similarity between a document and a query, which will tell us the relative importance of the query to the document.

I think a good corpus to use is the Harry Potter books since they are a good length and I find them interesting.

Dataset: - SORCERER STONE - CHAMBER OF SECRETS - PRISONER OF AZKABAN  
- GOBLET OF FIRE - ORDER OF THE PHOENIX  
- HALF-BLOOD PRINCE - DEATHLY HALLOWS (PART 1 AND 2)

### 1.2 Question/Proposal

How does the cosine similarity of each of the main character's names in all of the Harry Potter books relate to the amount of screen time they have in the Harry Potter movies?

### 1.3 Hypothesis

If the name of the character has a higher cosine similarity, then the weight of the name in the corpus is higher, meaning the importance of the person to the plot is greater, thus, they should have more screen time in the movie.

### 1.4 Zipf's Law

Before we begin using the model, let's explore why the model may make sense to use in the first place. What makes it work so well?

Zipf's law states that the rank of an entity - such as word - in a set will be inversely proportional to the number of times the entity shows up. For instance, in English, the most popular word - "the" - accounts for roughly 6% of all the words we say. By Zipf's law, the second most common word - "of" - will occur half as often as "the", which it mysteriously does, and so on for all subsequent words. This property seems to permeate all of nature and is even built into how we think. However, even though there have been many research papers studying Zipf's Law, no conclusive evidence has been discovered.

```
In [17]: HTML("""
    <div align="middle">
    
    </img></div>""")
    #Source = https://www.ahsthenest.com/archive/2018/03/19/zipfs-law/
```

```
Out[17]: <IPython.core.display.HTML object>
```

$$\text{number of occurrences} = \frac{1}{\text{rank}}$$

Moreover, this means that a majority of the words we write and say do not have much meaning. They are only there to space out and connect the meaningful/important concepts. This is why the Vector Space Model works so well. Even if a word appears hundreds of thousands of times in a corpus, it's importance will be diminished. In the Harry Potter books, the word "the" occurs all the time, but it gives no information about the plot of the books. The Vector Space Model takes advantage of the Zipf-ian nature of language to extract important terms from a corpus. Below, we use the Vector Space Model to do exactly that. In addition, we show that the Harry Potter books do in fact follow a Pareto distribution, which mathematically represent Zipf's law.

```
In [12]: #Frequency of terms in Sorcer of Stone
import re
from operator import itemgetter
from plotly import tools

datas=[]
layouts=[]
for i in range(7):
    file = open(f"B{i+1}.txt", 'rt')
    text = file.read()
    text=nlk.word_tokenize(text)
    frequency={}
    for word in text:
        if word.isalpha():
            count = frequency.get(word,0)
            frequency[word] = count + 1

    frequency = sorted(frequency.items(), key=lambda kv: kv[1], reverse=True)

    words = [item[0] for item in frequency]
    times = [item[1] for item in frequency]
    data = go.Bar(
        x=words[0:100],
        y=times[0:100]
    )

    datas.append(data)
    #layouts.append(layouts)
fig = tools.make_subplots(rows=4, cols=2,
```

```

specs=[[{}], {}],[{}], {}],[{}], {}], [{'colspan': 2}, None]],
subplot_titles=(books[0], books[1], books[2],
books[3],books[4],books[5],books[6]))

fig.append_trace(datas[0],1,1)
fig.append_trace(datas[1],1,2)
fig.append_trace(datas[2],2,1)
fig.append_trace(datas[3],2,2)
fig.append_trace(datas[4],3,1)
fig.append_trace(datas[5],3,2)
fig.append_trace(datas[6],4,1)

for i in range(len(books)):
    fig['layout']['xaxis'+str(i+1)].update(title='Word')

for j in range(len(books)):
    fig['layout']['yaxis'+str(j+1)].update(title='Frequency')

fig['layout'].update(showlegend= False)

fig['layout'].update(height=1000, width=1000,
                      title="Zip's law in all Harry Potter books")
py.iplot(fig, filename='basic-bar')

```

This is the format of your plot grid:

```

[ (1,1) x1,y1 ] [ (1,2) x2,y2 ]
[ (2,1) x3,y3 ] [ (2,2) x4,y4 ]
[ (3,1) x5,y5 ] [ (3,2) x6,y6 ]
[ (4,1) x7,y7   -       ]

```

**Out[12]:** <plotly.tools.PlotlyDisplay object>

## 1.5 Building Blocks

First, let's explore the building blocks that will allow us to determine the importance of a word/phrase to a corpus using the Vector Space Model.

### 1.5.1 Term Frequency

Term frequency (TF) is the number of times a word appears in a given corpus.

For example, the term frequency of "dope" in "That was dope." is 1.

### 1.5.2 Inverse Document Frequency

Inverse document frequency (IDF) is a measure of the rareness of a word in a corpus. It is defined as:

$$2 \quad idf_t = \log\left(\frac{N}{t_f}\right)$$

where N is the number of documents and  $t_f$  is the number of documents in the corpus which the term is present in.

The log function looks like:

```
In [2]: from IPython.display import HTML
HTML("""
<div align="middle">
<video width="100%" autoplay loop>
    <source src="VSM.mp4" type="video/mp4">
</video></div>""")
#See attached video VSM.mp4 if unable to view
```

```
Out[2]: <IPython.core.display.HTML object>
```

From the above graph, we can see that when the term isn't in that many documents, the input to the log function will be high, since N is constant, making the output higher.

### 2.0.1 TF-IDF

TF-IDF is the term frequency multiplied by the inverse document frequency. This is a measure of the uniqueness of a word in a document that exists in a corpus.

$$tf-idf = tf * idf$$

### 2.0.2 Cosine Similarity

The tf-idf values are then put into a vector that represents the word or document in a vector space. The cosine similarity is defined as:

$$CosineSimilarity = \frac{a \cdot b}{\|a\| \|b\|} \text{ where } a \text{ and } b \text{ are two document or word/pharse vectors}$$

```
In [3]: from IPython.display import HTML
HTML("""
<div align="middle">
<video width="100%" autoplay loop>
    <source src="CS.mp4" type="video/mp4">
</video></div>""")
#See attached video CS.mp4 if unable to view
```

```
Out[3]: <IPython.core.display.HTML object>
```

The higher the cosine similarity, the greater the dot product, since they are more in the same direction i.e. the projection of one vector onto the other is large. Each vector is divided by its magnitude to normalize the distance since we only care about the angle between them, so a unit vector is created in the direction of each vector. The cosine similarity is equal to  $\cos(\theta)$ , shown in the example above.

## 2.1 Data Exploration

Text files for books was obtained from <http://www.glozman.com/textpages.html>.

```
In [4]: from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.metrics.pairwise import cosine_similarity
        import nltk
        import pandas as pd
```

```
file = open("B1.txt", 'rt')
text = file.read()
text=nltk.word_tokenize(text)
print(text[0:10])
```

```
['Harry', 'Potter', 'and', 'the', 'Sorcerer', "'s", 'Stone', 'CHAPTER', 'ONE', 'THE']
```

The data looks like it is being tokenized properly. Now, we should find the tf-idf vectors, find the cosine similarity between each of the queries and books, get the screen times of each character in each movie, and see if there is a correlation between the two.

## 2.2 Data Analysis

```
In [5]: corpus=[]
        books = [
            "SORCERER STONE",
            "CHAMBER OF SECRETS",
            "PRISONER OF AZKABAN",
            "GOBLET OF FIRE",
            "ORDER OF THE PHOENIX",
            "HALF-BLOOD PRINCE",
            "DEATHLY HALLOWS (PART 1 AND 2)"
        ]
        num_books = 7
        for i in range(num_books):
            file = open(f"B{i+1}.txt", 'rt')
            text = file.read()
            corpus.append(text)

        characters=[
            "HARRY POTTER",
            "RON WEASLEY",
            "HERMIONE GRANGER",
            "VOLDEMORT",
            "DUMBLEDORE",
            "HAGRID",
            "DRACO MALFOY",
            "SIRIUS BLACK",
            "NEVILLE LONGBOTTOM"
```

```
]
```

```
cosine_similarities={}
for character in characters:
    corpus.append(character)
    vectorizer = TfidfVectorizer()
    tfidf_matrix = vectorizer.fit_transform(corpus)
    character_sim=[]
    for i in range(num_books):
        similarity = float(cosine_similarity(tfidf_matrix[i],tfidf_matrix[-1]))*1000
        sim_rounded = round(similarity,3)
        character_sim.append(sim_rounded)
    cosine_similarities[character]=character_sim

pd.DataFrame(cosine_similarities, index=books)
```

Out [5]:

	HARRY POTTER	RON WEASLEY	HERMIONE GRANGER	\
SORCERER STONE	141.351	47.602	29.664	
CHAMBER OF SECRETS	159.369	78.557	31.224	
PRISONER OF AZKABAN	149.567	68.196	44.165	
GOBLET OF FIRE	103.896	91.440	29.137	
ORDER OF THE PHOENIX	128.195	63.199	39.236	
HALF-BLOOD PRINCE	130.245	49.585	32.546	
DEATHLY HALLOWS (PART 1 AND 2)	124.315	60.461	54.717	

  

	VOLDEMORT	DUMBLEDORE	HAGRID	DRACO MALFOY	\
SORCERER STONE	5.436	22.853	53.403	13.619	
CHAMBER OF SECRETS	3.247	20.357	21.093	23.424	
PRISONER OF AZKABAN	3.532	11.281	27.362	16.343	
GOBLET OF FIRE	6.507	3.149	1.459	8.282	
ORDER OF THE PHOENIX	5.911	28.657	16.230	6.455	
HALF-BLOOD PRINCE	15.609	66.981	15.092	20.805	
DEATHLY HALLOWS (PART 1 AND 2)	16.536	33.336	11.902	1.443	

  

	SIRIUS BLACK	NEVILLE LONGBOTTOM
SORCERER STONE	5.867	12.972
CHAMBER OF SECRETS	5.478	3.831
PRISONER OF AZKABAN	31.253	11.076
GOBLET OF FIRE	10.485	1.214
ORDER OF THE PHOENIX	24.405	7.192
HALF-BLOOD PRINCE	7.490	3.825
DEATHLY HALLOWS (PART 1 AND 2)	7.959	0.435

\*the cosine similarity was multiplied by 1000 for easier viewing and distinction between the values. All calculations past this point will be based off the multiplied cosine similarity values.

## 2.3 Get screen times

Using IMDB's movie database (<https://www.imdb.com/list/ls027460372/>), we can get the screen times for each of the actors in the films.

```
In [6]: #screen times are in order of the books as shown above
screen_time={"HARRY POTTER": [72.75, 83.75, 74.5, 63.25, 61.75, 67.00, round(58+8/60, 2)],
             "RON WEASLEY": [28.25, 38.25, 21.25, 20.5, 21.0, 21.75, round(30+23/60, 2)],
             "HERMIONE GRANGER": [23.25, 15.5, 34.75, 16.5, 23, 20, 36],
             "VOLDEMORT": [2, 6.75, 0, 6.5, 2.25, 4.25, 7.75],
             "DUMBLEDORE": [9.75, 10.75, 6.25, 14.25, 7.5, 22.25, 3.25],
             "HAGRID": [16.5, 8.25, 5.25, 3.75, 2.75, 4, round(2+38/60, 2)],
             "DRACO MALFOY": [4.25, 7, 4, 2.25, 1.25, 8.25, 2.33],
             "SIRIUS BLACK": [0, 0, 11.5, 1.25, 7.5, 0, 0],
             "NEVILLE LONGBOTTOM": [3.25, 1.25, 3, 4.25, 9, 1.5, round(7/60, 2)]
            }
pd.DataFrame(screen_time, index=books)
```

```
Out [6]:
```

	HARRY POTTER	RON WEASLEY	HERMIONE GRANGER	\
SORCERER STONE	72.75	28.25	23.25	
CHAMBER OF SECRETS	83.75	38.25	15.50	
PRISONER OF AZKABAN	74.50	21.25	34.75	
GOBLET OF FIRE	63.25	20.50	16.50	
ORDER OF THE PHOENIX	61.75	21.00	23.00	
HALF-BLOOD PRINCE	67.00	21.75	20.00	
DEATHLY HALLOWS (PART 1 AND 2)	58.13	30.38	36.00	

  

	VOLDEMORT	DUMBLEDORE	HAGRID	DRACO MALFOY	\
SORCERER STONE	2.00	9.75	16.50	4.25	
CHAMBER OF SECRETS	6.75	10.75	8.25	7.00	
PRISONER OF AZKABAN	0.00	6.25	5.25	4.00	
GOBLET OF FIRE	6.50	14.25	3.75	2.25	
ORDER OF THE PHOENIX	2.25	7.50	2.75	1.25	
HALF-BLOOD PRINCE	4.25	22.25	4.00	8.25	
DEATHLY HALLOWS (PART 1 AND 2)	7.75	3.25	2.63	2.33	

  

	SIRIUS BLACK	NEVILLE LONGBOTTOM
SORCERER STONE	0.00	3.25
CHAMBER OF SECRETS	0.00	1.25
PRISONER OF AZKABAN	11.50	3.00
GOBLET OF FIRE	1.25	4.25
ORDER OF THE PHOENIX	7.50	9.00
HALF-BLOOD PRINCE	0.00	1.50
DEATHLY HALLOWS (PART 1 AND 2)	0.00	0.12

## 2.4 Get average cosine similarity and screen time for each character

```
In [8]: average_similarities = {}
        average_screen_times = {}
```

```

for character in characters:
    average_similarities[character]=sum(cosine_similarities[character])\
    /len(cosine_similarities[character])
    average_screen_times[character]=sum(screen_time[character])\
    /len(screen_time[character])

HTML("""
<div align="middle">

</img></div>""")

```

Out[8]: <IPython.core.display.HTML object>

## 2.5 Plot on graph to see relationship

```

In [11]: import warnings
warnings.filterwarnings('ignore')
import plotly.plotly as py
import plotly.graph_objs as go
from scipy import stats
import numpy as np

sims = np.array(list((average_similarities.values())))
times = np.array(list((average_screen_times.values())))

slope, intercept, r_value, p_value, std_err = stats.linregress(sims,times)
line = slope*sims+intercept

trace1 = go.Scatter(
    x = sims,
    y = times,
    mode = 'markers',
    text = characters
)

trace2 = trace2 = go.Scatter(
    x=sims,
    y=line,
    mode='lines',
    marker=go.Marker(color='rgb(31, 119, 180)'),
    name='Fit'
)

layout= go.Layout(
    title= 'Average screen time vs cosine similarity \
for harry potter characters over all books and movies',
    hovermode= 'closest',
    xaxis= dict(

```



```

        title= 'Cosine Similarity',
        ticklen= 5,
        zeroline= False,
        gridwidth= 2,
    ),
    yaxis=dict(
        title= 'Screen Time (minutes)',
        ticklen= 5,
        gridwidth= 2,
    ),
    showlegend= False
)
data = [trace1, trace2]
fig= go.Figure(data=data, layout=layout)

# Plot and embed in ipython notebook!
py.iplot(fig, filename='basic-scatter')

```

Out[11]: <plotly.tools.PlotlyDisplay object>

As we can see from the graph above, there is a strong positive correlation between cosine similarity and screen time. (more discussed in conclusion)

## 2.6 Levels of Linguistic Analysis

### 2.6.1 Syntax

Word and phrase creation is a difficult task. To make it easier on ourselves, we formulate phrases through a combination of meaningful words and connecting words. This gives us time to think without pausing too often. Beniot Mandelbrot, a famous mathematician, said that nature always looks for the path of least energy and language may have evolved in a similar process in order to keep language as fluid and understandable as possible.

### 2.6.2 Semantics

As explored above, the meaning of speech and writing is held in a select few words. These words don't occur often, but that is what makes them influential. The downside to this is that meaning takes long to convey. The Zipf-ian nature of language seems to have caused verbosity to be an inherent part of our language. If every word we said had equally influential meaning, then communication could occur much more efficiently. The question is whether the human brain would be capable of producing coherent information that quickly. Unfortunately, that is probably not the case. Moreover, for the listener, the meaning is not always easy to find since it exists in a sea of meaningless words. You have to be constantly listening for the important parts because if you lose focus for a brief time, you may miss them. In all, it would make sense for language to evolve for ease of production and retention, but that doesn't seem to be the case. It may be that we are in the nascent stages of language and it may take longer for it to reach its maximum potential.

## 2.7 Conclusion

In this experiment, I was attempting to uncover the interlinked nature of literature and film through analyzing the content of the Harry Potter books. More specifically, I wanted to see how well the characters in the books were depicted on-screen. To do this, I used the cosine similarity of a character's name and each book as a proxy for the relative importance of the term. I believe this is a fair assumption because, many times, the use of a name indicates that the character had a line of dialogue, which directly relates to them being onscreen. In addition, the use of a character's name in any regard works to increase their importance in the overall scheme of the book as it shows they are more integral to the plot. My hypothesis was that the higher the cosine similarity (loosely translated as the importance of the name in the book), the higher the screen time for that character.

This experiment confirmed my hypothesis and showed that the importance of the characters in the books was accurately portrayed through the film. As seen from the graph above, the relationship is linear with an  $r^2$  value of 0.974, which indicates a strong positive correlation. This implies that the movie closely mirrored the events in the books and each character's contributions. Interestingly, as much as Voldemort is a major component of the plot of many of the films, his screen time doesn't correlate that well with the weight of his name in each of the books. This makes sense though as Voldemort is talked about a lot by other characters, increasing his cosine-similarity, but only present during the tension and action filled scenes towards the end of films.

In all, we explored how Zipf's law characterizes language, allowing for the Vector Space Model to be used effectively, and how different expressions of language - books and movies - relate to one another.

## 3 Appendix

In [ ]: *#using manim math engine, code for visualization of log function*

```
class LOG(GraphScene):
    CONFIG = {
        "x_min" : -5,
        "x_max" : 5,
        "y_min" : 0,
        "y_max" : 25,
        "graph_origin" : DOWN*3,
        "function_color" : RED,
        "axes_color" : GREEN,
        "x_axis_label": None,
        "y_axis_label": None,
    }

    def construct(self):
        self.setup_axes(animate=True)
        graph = self.get_graph(lambda t : np.exp(t), color=self.function_color)
        self.play(ShowCreation(graph))
```

In [ ]: *#using manim math engine, code for visualization of cosine similarity*

```
class CS(Scene):
    CONFIG = {
```

```

"x_min" : -10,
"x_max" : 10,
"y_min" : -10,
"y_max" : 10,
"function_color" : RED,
"axes_color" : GREEN,
}

def construct(self):
    plane = NumberPlane()
    vector1=Vector(3*UP+RIGHT)
    vector2=Vector(3*RIGHT+UP)
    group=VGroup(vector1,vector2)
    x_line=Line(ORIGIN,RIGHT, color=RED)
    y_line=Line(RIGHT,3*UP+RIGHT, color=GREEN)
    ordered_pair = TexMobject("[1,3]")
    ordered_pair.next_to(vector1.get_tip())
    self.play(ShowCreation(
        plane,
        submobject_mode = "lagged_start",
        run_time = 3
    ))
    self.play(GrowFromPoint(
        group,
        ORIGIN,
        submobject_mode = "lagged_start"),
        run_time=3)
    self.wait()
    self.play(ShowCreation(
        x_line
    ))
    self.play(ShowCreation(
        y_line
    ))
    self.play(Write(
        ordered_pair
    ))
    self.play(FadeOut(
        x_line
    ),FadeOut(
        y_line
    ))
    x_line=Line(ORIGIN,3*RIGHT, color=RED)
    y_line=Line(3*RIGHT,3*RIGHT+UP, color=GREEN)
    ordered_pair_2 = TexMobject("[3,1]")
    ordered_pair_2.next_to(vector2.get_tip())
    self.play(ShowCreation(
        x_line

```

```

))
self.play(ShowCreation(
    y_line
))
self.play(Write(
    ordered_pair_2
))
self.play(FadeOut(
    x_line
), FadeOut(
    y_line
))
theta = Arc(np.arctan(3)-np.arctan(1/3), start_angle=np.arctan(1/3))
angle = TexMobject("\\theta")
angle.next_to(theta)
angle.shift(UP*.3+LEFT*.3)
equation = TexMobject("cos(\\theta)=\\frac{a\\cdot b}{|a| \\times |b|}")
equation.shift(2.5*LEFT+2*DOWN)
equation_2 = TexMobject("cos(\\theta)=\\frac{[1,3]\\cdot \\[3,1]}{|[1,3]| \\times |[3,1]|}")
equation_2.shift(2.5*LEFT+2*DOWN)
equation_3 = TexMobject("cos(\\theta)=\\frac{1\\times 3 + 3 \\times 1}{\\sqrt{1^2+3^2} \\times \\sqrt{3^2+1^2}}")
equation_3.shift(2.5*LEFT+2*DOWN)
equation_4 = TexMobject("cos(\\theta)=\\frac{6}{\\sqrt{10} \\times \\sqrt{10}}")
equation_4.shift(2.5*LEFT+2*DOWN)
equation_5 = TexMobject("cos(\\theta)=\\frac{3}{5}=Cosine \\hspace{1.5mm} Simi")
equation_5.shift(2.5*LEFT+2*DOWN)
self.play(
    ShowCreation(theta),
    ShowCreation(angle)
)
self.play(Write(equation))
self.wait()
self.play(FocusOn(ordered_pair), FocusOn(ordered_pair_2))
self.wait()
self.play(FadeOut(equation), Write(equation_2), run_time=2)
self.wait()
self.play(FadeOut(equation_2), Write(equation_3), run_time=2)
self.wait()
self.play(FadeOut(equation_3), Write(equation_4), run_time=2)
self.wait()
self.play(FadeOut(equation_4), Write(equation_5), run_time=2)
self.wait()

```